

Digital Scholarship in the Humanities Suitability and Revision Plan for litpaper

Executive summary

The `MuhammadTahaBinZaeem/litpaper` repository is **much stronger as an empirical package than as a DSH-ready paper**. Its biggest assets are unusually good pipeline traceability, deterministic corpus construction, author/title blinding in the rewrite prompts, explicit QC, passage-level split discipline that prevents original/rewrite leakage, and a reproducibility mindset backed by checker scripts and GitHub Actions. The core result is also real rather than cosmetic: on the repo's own nearest-centroid test setup, author-classification macro-F1 drops from **0.7769** on originals to **0.3958** after paraphrase, **0.5607** after modernisation, and **0.4636** after simplification; inter-author distance likewise contracts under rewriting; and the repo already includes bootstrap intervals showing those losses are positive on the test split. That is publishable raw material. ¹

The problem is that **DSH will not reward pipeline bravado on its own**. The journal explicitly says that AI/data-science work is out of scope if it lacks theoretical relevance or direct application to Digital Humanities, and it does not publish mere project or method descriptions. DSH also requires a structured abstract, a data availability statement, and an AI Disclosure Statement. So the repo's current state is a good **research object**, but not yet a good **DSH argument**. If you submit this as "we benchmarked one LLM and style got flatter", you are inviting a desk rejection as a generic AI/stylometry paper. If you recast it as a **digital literary study of how controlled rewriting changes the measurability, transferability, and interpretability of authorial style**, and add the right robustness checks and humanities-facing interpretation, the paper becomes credible for DSH. ²

My harsh assessment is this: **current DSH fit is 5.8/10**. The repo shows serious work, but the likely editorial reaction is "interesting, but is this a DH paper or a narrow AI-stylometry experiment?" With the concrete upgrades below, I think you can move that to roughly **7.5–8.0/10 fit**, which is the zone where a send-out for peer review becomes plausible rather than optimistic. That is not the same thing as guaranteed acceptance; it means the paper stops looking like it wandered in from NLP and starts looking like it belongs in DSH. ³

Repository audit against DSH expectations

The repo's strongest and weakest features are easiest to see side by side.


Aspect	Evidence in the repo	Why DSH reviewers will care	Verdict
Corpus balance and selection discipline	<code>logs/step_08_status.md</code> reports a deterministic selection of 360 original passages across 6 authors and 12 works , with 60 passages per author and 30 per work ; <code>metadata/selection_summary.csv</code> confirms deterministic even spread and a 550–645 word realised range. ⁴	Strong anti-bias design; reads as deliberate humanities corpus construction rather than ad hoc scraping.	Strong
Rewrite protocol validity	<code>docs/rewrite_protocol.md</code> locks three conditions (<code>paraphrase</code> , <code>modernize</code> , <code>simplify</code>), forbids summarisation and author/title leakage, and requires prompt hashes, source-text hashes, rewrite hashes, model/version recording, QC flags, and seeded generation where available. ⁵	This is exactly the kind of methods transparency that makes stylometric claims more credible.	Strong
Split hygiene	<code>logs/step_13_modeling_matrix_report.md</code> and the audit state that all versions of a passage stay in the same split, preventing leakage across original and rewritten forms. ⁶	Excellent. This is one of the first things a competent reviewer will check.	Strong
Reproducibility auditing	<code>logs/final_reproducibility_audit_report.md</code> shows all major checker scripts passing and 36 critical files checked, 0 missing ; workflows automate source migration and later analysis stages. ⁷	DSH is increasingly receptive to rigorous digital scholarship when it is auditable and embodied in research software.	Strong
Empirical signal strength	Step 14–19 reports show a meaningful original baseline, degradation under rewriting, residual same-condition signal, inter-author distance contraction, feature-family vulnerability analysis, and bootstrap intervals. ⁸	This is enough for a real paper, not merely a pilot note.	Strong

Aspect	Evidence in the repo	Why DSH reviewers will care	Verdict
Generator narrowness	The repo's own audits admit a single Gemini model and frame the supported claim narrowly as one controlled rewriting setup. ⁹	This is the single biggest external-validity weakness. Reviewers will ask whether you have discovered an LLM effect or a Gemini effect.	Weak
QC retention risk	<code>logs/step_10_rewrite_generation_report.md</code> records 111 warning rows, 44 hard length warnings , and 1 JSON parse failure handled using raw content; Step 11 carries those warnings forward with no master failures. ¹⁰	Not fatal, but absolutely reviewable. If you leave this untreated in prose, it looks like soft contamination of the main experiment.	Weak
Corpus generalisability	The audit explicitly says the package is limited to six public-domain fiction authors, English literary prose only , and passage-level rather than full-book attribution. ¹¹	DSH will tolerate a bounded corpus if you frame it honestly, but it will punish universal conclusions.	Weak
Full-source reproducibility gap	Step 4 states that the cleaned full corpus was generated locally and checked, but the cleaned source texts are not fully committed in the public repo because of source/package constraints. ¹²	This is not devastating because the selected passage texts are committed, but it is still a vulnerability for full upstream reproduction.	Weak


Aspect	Evidence in the repo	Why DSH reviewers will care	Verdict
Environment specification	<code>environment.yml</code> is serviceable, but <code>requirements.txt</code> is still explicitly labelled a dependency stub rather than a final pinned environment. ¹³	DSH reviewers will not love sloppy dependency management in a paper that makes strong reproducibility claims.	Weak

Two strengths deserve special emphasis. First, the repo has already done the hard part that many DH papers dodge: it has created a **traceable chain from corpus to claims**, including a claim-support matrix and final audit. Second, the paper’s main result is not a trivial “AI changed the text” observation; the repo shows both **classification degradation** and **inter-author distance contraction**, which is a better warrant for talking about style flattening than raw classifier scores alone. ¹⁴

The weaknesses are also not cosmetic. The current design still leaves an obvious reviewer opening: **author is partly confounded with work/topic** because each author contributes only two works, and the main split appears to be passage-level rather than held-out-work evaluation. Stylometry has long been vulnerable to content leakage and within-author variation problems; classical DSH stylometry work explicitly warns that differences within an author’s own corpus matter, while more recent NLP work on content-independent style representation shows that author signals can be entangled with topic unless you control for them aggressively. As it stands, a sceptical reviewer can say: “You may be measuring loss of work-specific lexical signatures as much as loss of authorial style.” That criticism would hurt, because right now it would not be entirely wrong. ¹⁵



The repo's own numbers make the argumentative hierarchy clear: **paraphrase is the most damaging condition for cross-condition transfer**, not simplify, and that is intellectually interesting. It means the paper should stop implying a simple "more rewriting severity equals more style destruction" story. Your data support something subtler: **light lexical-syntactic substitution may confuse models trained on original texts more than heavier readability-oriented rewriting does**, perhaps because simplification leaves behind stable author-conditioned constraints that are easier to relearn within-condition. That is a much better DSH discussion than a blunt flattening slogan. ¹⁶



Distance contraction strengthens the case because it is not tied to any one classifier. On the all-feature measure, mean inter-author distance falls from **0.4665** in the original condition to **0.3822** for paraphrase, **0.4171** for modernise, and **0.4053** for simplify; the function-word-only analysis points in the same direction. That gives you a stronger humanities claim: controlled rewriting does not merely lower model performance, it **compresses the stylistic space between authors**. ¹⁷

DSH fit and likely editorial reaction

DSH is open to methodological, experimental, and applied DH research, but it is also explicit that it will not publish **generic AI/data-science work** without clear theoretical or practical relevance for DH. At the same time, the journal is plainly publishing AI-related work right now: recent advance articles include work on AI-generated art, AI-assisted hermeneutics, AI and virtual archives, AI in cultural heritage, LLM-based story generation, AI-driven translation of Buddhist scriptures, and reception experiments on AI-versus human-authored poetry. So the answer is not “DSH hates AI”. The answer is “DSH will publish AI when it is unmistakably a humanities paper.” ¹⁸

My fit judgement is therefore:

Item	Assessment
Current DSH fit	5.8/10
If submitted now, likely editorial view	“Potentially interesting, but under-theorised and too close to a single-model AI benchmark.”
Main desk-reject risk	Out-of-scope AI/data-science framing without enough DH contribution and literary interpretation.

Item	Assessment
Main peer-review risk if sent out	Topic/style confound, single-LLM generalisability, insufficient humanities payoff, and warning-row robustness.
Revised fit after strong changes below	7.5–8.0/10

That score is a judgement call, but it is grounded in the mismatch between the repo's current strengths and DSH's stated expectations. The journal wants a structured abstract with a declared contribution to DH, requires a Data Availability Statement, requires an AI Disclosure Statement when AI tools are used, and expects a paper rather than a project log. Your repo already serves the paper; now the paper has to serve the journal. ¹⁹

If I were predicting the first editorial reaction, I would split it into two scenarios. **Scenario one:** the title and abstract sound like generic LLM evaluation. In that case, I think the editor is likely to conclude that the paper belongs elsewhere, because DSH's scope warning is explicit. **Scenario two:** the title and abstract foreground a DH problem — measurable authorial style, controlled rewriting as intervention, implications for digital literary method, and reproducible analysis of a literary corpus. In that case, I think the editor could reasonably send it to review, especially because DSH is presently publishing AI-centred work when it is clearly integrated with humanities questions. That is the entire game here: same data, radically different editorial destiny depending on framing. ¹⁸

There is also an administrative point worth stating plainly because it will matter at submission time. DSH full papers should normally remain within **9,000 words**, and the abstract must be structured and no longer than **250 words**. The journal offers both standard-licence and open-access publication; open access can involve a charge unless your institution is covered by an Oxford read-and-publish agreement, whereas the standard route does not require you to pay an OA fee merely to publish. ²⁰

Concrete manuscript redesign

The paper needs a redesign, not a polish. The redesign has three principles.

First, **move the centre of gravity from “an LLM changed some stylometric metrics” to “controlled rewriting reveals something about the stability, transferability, and measurement of literary style.”** DSH will respect the latter. The former sounds like a light NLP paper with nineteenth-century garnish. This is not semantics. It changes what counts as literature review, what counts as methods, and what counts as contribution. ²¹

Second, **stop presenting the repo as the contribution.** DSH says it does not publish project descriptions or bare methodologies. The contribution must instead be framed as a humanities argument supported by a reproducible digital method. That means the manuscript has to interpret why paraphrase collapses transfer most, why rewritten texts still retain same-condition author signal, what feature families are most vulnerable, and what this means for literary style as an object of digital scholarship. ²²

Third, **add visible humanities interpretation.** Right now the repo proves that the empirical package is ready for writing, but it also says the literature framing and interpretation still need to be written. DSH reviewers will not be impressed by a paper that never returns from metric space to literary prose. You need close-reading case studies that show what the measured contractions actually look like in Austen, Poe, Twain, and so on. Otherwise the article remains method-forward and interpretively thin. ²³

Recommended paper architecture

A strong DSH full paper here would look like this, with a disciplined word budget.

Section	Target words	What must be in it
Introduction	1,000–1,200	Literary style as a DH object; why controlled rewriting matters; the specific gap in LLM–stylometry work; your research questions.
Related work	1,000–1,200	Classical stylometry, DSH authorship studies, recent DSH AI papers, translation/stylometry, and content-vs-style confounds.
Corpus and rewrite protocol	1,000–1,200	Source selection, passage balancing, prompt blinding, QC, and why your interventions count as controlled rewriting rather than summarisation or imitation.
Stylometric representation and experiments	1,100–1,300	Feature families, split design, baseline models, transfer tests, same-condition tests, distance measures, bootstrap design.
Results	1,600–1,900	The actual outcomes, with one subsection each for baseline, degradation, residual signal, distance contraction, feature families, and robustness.
Discussion	1,100–1,300	Literary and DH implications; what “flattening” means and does not mean; interpretive case studies; limits.
Conclusion and end matter	400–600	Narrow claim, future work, data availability, AI disclosure.

That structure fits the journal’s normal full-paper limit and gives the methods enough room without suffocating the interpretive discussion. ²⁴

Non-generic edits you should make

You should make these edits literally, not in spirit.

Manuscript location	Current problem	Exact fix
Title	Too easy to overclaim “LLMs flatten style”.	Replace with a title that names controlled rewriting , authorial signal , and the bounded corpus/model setup .
Opening paragraph	Likely to start with AI hype.	Start instead from authorial style as a long-standing problem in computational literary study , then introduce rewriting as an intervention on measurability.

Manuscript location	Current problem	Exact fix
End of introduction	Probably lacks explicit DH contribution.	Add a final paragraph beginning: “This article contributes to Digital Humanities by...” and name corpus design, reproducibility, and implications for computational literary analysis.
Methods	Pipeline-heavy but possibly too repo-like.	Cut procedural noise; keep only what a sceptical reviewer needs to understand validity, leakage control, and reproducibility.
Results	Risk of becoming a metric dump.	Each metrics subsection must end with a one-sentence humanities interpretation, not just a numerical summary.
Discussion	Risk of generic limitations.	Add a dedicated “What this does not show” subsection: not all LLMs, not all literatures, not all notions of style, not human indistinguishability.
End matter	Missing DSH-required statements.	Add Data availability and AI Disclosure Statement in DSH format.

Everything above comes directly from the gap between the repo’s readiness and DSH’s requirements. The journal mandates the abstract structure and end-matter statements; the repo’s own audit says literature framing, final title wording, and journal formatting still remain. ²⁵

Sample structured abstract for DSH

Purpose: This article investigates how controlled large-language-model rewriting affects the computational separability of literary authorial style. Rather than treating rewriting as a generic AI benchmark, it approaches paraphrase, modernisation, and simplification as structured interventions on the measurability of prose style within digital literary study. **Design/methodology/approach:** I construct a balanced corpus of 360 passages from 12 public-domain works by 6 nineteenth-century authors and generate 1,080 blinded rewrites under three controlled conditions. I then model authorial separability using 205 stylometric features, train-on-original transfer tests, same-condition classification, Burrows’ Delta distance analysis, feature-family comparisons, and paired bootstrap confidence intervals. **Findings:** Authorial signal is clearly measurable in the original passages, but train-on-original transfer degrades after rewriting. The largest loss occurs under paraphrase, while modernisation and simplification also contract inter-author stylometric distance. Rewritten texts nevertheless retain residual author-specific signal when analysed within-condition, showing that rewriting weakens transferability more than it erases stylistic structure. **Originality:** The study combines a controlled rewrite protocol, explicit leakage control, and multiple stylometric analyses to examine LLM rewriting as a methodological problem for computational literary analysis. **Contribution to the field of Digital Humanities:** The article offers a reproducible case study of how generative rewriting changes the measurement of literary style and clarifies what kinds of claims digital humanists can and cannot make from stylometric evidence after AI-mediated textual transformation. ²⁶

Sample introduction paragraph

This article treats LLM rewriting not primarily as a problem of model capability, but as a problem of **digital literary method**. Stylometry has long relied on the hypothesis that authorial style leaves measurable traces across texts, yet recent work in DSH shows both that modern generative systems can

imitate aspects of literary style and that machine-mediated transformation can disturb stylometric attribution in non-trivial ways. What remains underexplored is what happens when rewriting is made controlled, blinded, and comparable across literary passages: does authorial signal simply disappear, or does it become less transferable while still surviving within the rewritten condition itself? By answering that question on a balanced public-domain corpus, the article shifts the discussion from AI spectacle to the stability of literary style as an object of digital scholarship. ²⁷

Sample methods paragraph

To isolate rewriting from obvious confounds, the study uses a deterministic selection of 360 original passages from 12 works by 6 authors, balanced by author and work and filtered for length and sentence count. Rewriting was performed under a locked protocol with three non-summarising conditions — paraphrase, modernise, and simplify — while the generator was blinded to author names and titles. Each rewritten row records model metadata, prompt hashes, source and output hashes, length ratios, and QC status. All versions of the same passage remain in the same train/validation/test split, preventing content leakage between original and rewritten forms. The analysis then combines original-condition baselines, train-on-original transfer experiments, same-condition classification, inter-author distance measures, feature-family comparisons, and paired bootstrap intervals. ²⁸

Sample results paragraph

Original-condition classification confirms that authorial separability is measurable in the corpus before rewriting. For the nearest-centroid model, macro-F1 on the original test split is 0.7769; when the same model is transferred to rewritten texts, performance falls to 0.3958 for paraphrase, 0.5607 for modernisation, and 0.4636 for simplification. These losses remain positive under paired bootstrap confidence intervals. At the same time, rewritten texts are not stylistically homogeneous: same-condition classification shows survival ratios of 0.8030 for paraphrase and 0.8782 for both modernisation and simplification, while inter-author distance also contracts under all three rewrite conditions. Read together, these results suggest that controlled rewriting compresses stylometric distance and weakens cross-condition transfer without eliminating recoverable author signal altogether. ²⁹

Sample discussion paragraph

The evidence supports a narrow but significant DH claim. Under one controlled Gemini rewriting setup, authorial style becomes less transferable from original to rewritten prose, but not fully anonymous. This matters for computational literary studies because it cautions against both naïve continuity and naïve collapse: rewritten texts should not be assumed either to preserve original stylistic relations intact or to erase them completely. More importantly, the strongest degradation under paraphrase shows that apparent semantic fidelity can coincide with substantial movement in stylometric space. For digital literary scholarship, that means AI-mediated textual transformation is not a neutral preprocessing step; it is an intervention that changes what style can be measured and how confidently authorship-like inferences can be carried across textual states. ³⁰

Experiments, robustness checks, and the reproducibility package DSH will expect

If you want to maximise publication odds, you need to add experiments that answer the reviewers before they ask.

Priority	Experiment	Exact design	Why it matters
Essential	Multi-LLM replication	Re-run the full 360-passage × 3-condition design on at least three models used in recent DSH-adjacent style-imitation work: GPT-4o , Gemini 1.5 Pro , and Claude Sonnet 3.5 ; keep prompts blinded, <code>temperature=0.2</code> , <code>top_p=1.0</code> , same length tolerance, and log seeds where supported; for models without deterministic seeding, do 3 runs/model and report variability.	Separates “Gemini effect” from “rewriting effect”. Recent DSH work on AI style imitation already treats multi-model comparison as normal. ³¹
Essential	Held-out-work validation	Because each author has two works, run leave-one-work-out evaluation: train on one work per author and test on the other, then reverse. Report macro-F1 and distance contraction.	This is the sharpest available answer to the “you are measuring topic/work, not author” objection. ³²
Essential	Warning-row sensitivity analysis	Recompute Step 15–19 metrics on three subsets: all rows , QC pass only , and pass + non-hard-length-warning .	Neutralises the QC-warning objection immediately. ³³
Essential	Semantic fidelity audit	Randomly sample 10% of rewritten passages per condition (108 rewrites total). Have 2 human annotators code added facts, omitted facts, narrative-order changes, and tone drift; report agreement and discrepancy rates.	Your current protocol promises semantic preservation; reviewers will want evidence that the intervention actually stayed controlled. ³⁴
Strongly recommended	Human reading study	72 participants , each reading 24 passages in a Latin-square design balanced across conditions and authors. Collect: readability, perceived stylistic distinctiveness, human/AI guess, and same-author/different-author similarity judgement. Analyse with mixed-effects models.	DSH is currently publishing AI reception work; adding reader evidence strengthens the humanities payoff enormously. ³⁵
Strongly recommended	Prompt/protocol ablation	Compare blinded prompts against a deliberately author-revealing version on a smaller subset, e.g. 60 passages .	Shows that your blinding choice matters methodologically, not just procedurally. ³⁴
Strongly recommended	Length-band robustness	Re-run transfer analysis in passage-length bands: 450–500 , 501–575 , 576–650 words .	Tests whether degradation is partly a length effect. ³⁶

Priority	Experiment	Exact design	Why it matters
Nice but valuable	Leave-one-feature-family-out ablation	Run nearest-centroid on full features minus one family at a time, in addition to family-only analyses.	You already know which families are vulnerable; this tells you which ones drive the full-system result. ³⁷

The repo already has a decent statistical spine, but you can make it substantially harder to dismiss with one extra layer of reporting. In addition to macro-F1 and bootstrap intervals, report: **distance ratios**, **survival ratios**, **paired effect sizes on per-passage nearest-author margin changes**, and **per-condition permutation tests** for the null that degradation is zero. Step 19 already uses **2,000 bootstrap iterations**; for the final paper, I would move that to **10,000** for the principal inferential table, especially if you are going to make a headline claim from the intervals. ³⁸

The reproducibility package you should submit or cite

DSH itself requires a Data Availability Statement, and DH debate increasingly treats reproducibility as an epistemic as well as technical issue. Your repo is already ahead of many papers, but the package should be tightened into a submission-grade bundle. ³⁹

A convincing package would contain the following.

Package element	Current repo state	What to improve before submission
Read-only archival snapshot	Strong GitHub repo with final audit. ⁴⁰	Create a versioned DOI snapshot of the submission-state repo and cite that in the paper.
Corpus release	Selected passage texts are committed; full cleaned corpus is not fully public. ⁴¹	In the paper, be explicit that the paper-level corpus is the selected-passage layer; provide those files cleanly, with checksums and provenance.
Environment capture	<code>environment.yml</code> exists; <code>requirements.txt</code> is still a stub. ¹³	Freeze exact versions; add a lockfile or container recipe.
Reproduction instructions	Checker scripts and workflows exist. ⁴²	Add a single top-level reproduction script and a plain-English <code>REPRODUCING_THE_PAPER.md</code> .
Results package	Tables, manifests, reports, and SVG figures exist. ³⁸	Add a manuscript-facing results appendix mapping each table/figure directly to the paper section that cites it.
AI use transparency	Technical logging is excellent. ⁵	Add a formal AI Disclosure Statement in the manuscript and cover letter.

Package element	Current repo state	What to improve before submission
Data transparency	Evidence exists, but no manuscript statement yet.	Add a formal Data availability statement with explicit files, archive locations, and access conditions.

A DSH-compliant end matter could look like this:

Data availability: “The submission’s selected-passage corpus, rewrite protocol, feature tables, modelling outputs, figures, and reproducibility scripts are available in the archived project snapshot cited in this article. The paper-level corpus used for the analyses consists of the selected passage layer and derived feature tables. Upstream full-source cleaning artefacts are documented through checksums, manifests, and scripts; where full source redistribution is constrained, regeneration instructions and provenance records are provided.” ⁴³

AI Disclosure Statement: “This manuscript was prepared with the assistance of generative AI tools for language refinement and drafting support. Large language models were also used as the object of study in the controlled rewriting experiments described in the Methods section. All model-generated content, including rewritten passages produced for analysis and any AI-assisted manuscript text, was checked and approved by the author. Model names, versions, prompt controls, and run metadata are reported in the project archive and Methods section.” That wording aligns with OUP’s requirement that AI use be declared with tool names, versions, purpose, and author verification. ⁴⁴

Anonymisation and review-safe submission steps

The public DSH/OUP material I reviewed does **not clearly state the journal’s exact peer-review model**, so I would treat anonymisation as a **risk-reduction norm**, not as a documented universal rule. OUP says review models vary by journal, and DSH’s public pages do not clearly spell out a review-blinding procedure in the material I found. At the same time, published DSH PDFs visibly retain placeholders like “[redacted for peer review]”, which strongly suggests anonymised handling in at least some DSH workflows. ⁴⁵


So do this anyway:

Step	What to do
Manuscript file	Remove author name, affiliation, ORCID, acknowledgements, grant numbers, personal GitHub handle, and self-identifying repository references.
Self-citations	Replace with neutral forms such as “Author” where required by the review setup, and restore the full references after acceptance.
Repo for review	Prepare a blinded mirror or archived package that strips your username from URLs, repository labels, PDF metadata, and figure metadata.
Supplementary files	Remove hidden metadata from Word, PDF, SVG, and CSV exports; check PDF properties and SVG <code><metadata></code> blocks.
Cover letter	State clearly that a blinded reproducibility package is available for review and that the named repository/DOI will be restored upon acceptance.

Step	What to do
Post-acceptance	Restore full links, acknowledgements, funding details, and repository identity.

Revision roadmap, reviewer objections, and exemplar models to emulate

The revision needs triage. Do not try to improve everything at once.



Priority checklist

Time horizon	Highest-value tasks
Next 24 hours	Rewrite the title and abstract around DH contribution; add Data availability and AI Disclosure end matter; run pass-only and pass-without-hard-length-warning sensitivity tables; write one paragraph explicitly acknowledging single-model and corpus-bounded limitations.
Within 1 week	Write the full literature review; add at least three close-reading mini-cases ; produce one held-out-work robustness analysis; freeze a versioned reproducibility bundle; finalise dependency pinning.
Pre-submission	Add multi-LLM replication if at all possible; update all figures for journal style; trim to $\leq 9,000$ words; check anonymisation; ensure every major claim points to a table/figure and a reproducible file.
Major work if time permits	Human reading study; semantic-fidelity annotation study; leave-one-family-out ablation; expanded corpus with more authors or non-fiction/later prose.

Likely reviewer objections and scripted rebuttals

Likely objection	Why it is dangerous	Rebuttal you should be able to give
"This is just generic AI evaluation, not DH."	This can kill the paper at desk stage.	"The article does not benchmark general model quality; it uses controlled rewriting as an intervention on a core DH object — measurable literary style — and ties the results to stylometry, corpus construction, and interpretation in computational literary studies."

Likely objection	Why it is dangerous	Rebuttal you should be able to give
"Your result is model-specific because you used one generator."	Fair criticism.	"We agree and therefore delimit the main claim to one controlled rewriting setup; in the revised version, we add cross-model replication / explicitly recast this as a bounded case study."
"You may be measuring topic or work, not author."	Also fair.	"We already prevent passage-version leakage through split discipline, and in revision we add leave-one-work-out evaluation to separate work-level lexical cues from cross-work authorial structure."
"Warning rows undermine the main results."	Easy reviewer hit.	"We reran all principal analyses on pass-only and non-hard-warning subsets; the direction and interpretation remain stable."
"There is not enough humanities interpretation here."	Very likely in DSH.	"We added close-reading case studies showing what the measured contractions look like in actual prose, and we discuss how the metrics change the interpretation of authorial style rather than treating them as self-explanatory."
"You overclaim that style is flattened or erased."	Dangerous because your own data reject total erasure.	"We now use the narrower claim that controlled rewriting contracts inter-author distance and weakens train-on-original transfer while leaving residual same-condition author signal."

Recommended framing language

Use phrases like these:

- **"under one controlled rewriting setup"**
- **"reduces the transferability of original-condition authorial signal"**
- **"contracts measurable inter-author stylometric distance"**
- **"does not eliminate residual author-specific structure"**
- **"within a balanced corpus of nineteenth-century English prose fiction"**
- **"a reproducible DH case study rather than a universal model claim"** ⁴⁶

Do **not** use phrases like:

- "LLMs erase authorial style"
- "AI can write indistinguishably like human authors"
- "this proves stylometry fails under AI"
- "these findings generalise to literary writing broadly"

Those are all bigger than your evidence. The repo's own audit already advises narrower language. ⁴⁷

Title options that are good enough for DSH

- **Controlled LLM Rewriting and the Contraction of Authorial Distance in Nineteenth-Century English Fiction**

- **How Much Authorial Signal Survives Controlled Rewriting? A Stylometric Study in Digital Literary Analysis**
- **From Original to Rewrite: Measuring the Transferability of Literary Style after Paraphrase, Modernisation, and Simplification**
- **Rewriting Literary Prose with Large Language Models and What It Changes in Stylometric Authorship Analysis**

Exemplar DSH papers and editorials to emulate

Two editorial or field-positioning pieces are worth imitating for tone and framing. The special-issue introduction on **collaboration as opportunity** is a good model for writing oneself into a DH conversation rather than merely presenting results, while the conceptual paper on **knowledge complexity and the five “apories”** is a strong model for explaining why a computational method matters epistemically, not only technically. ⁴⁸

For recent article models, the following are especially useful.

DSH article	Contribution	Methods	Approximate length	Data / supplement signal	Why it worked as DSH
George Mikros, “Beyond the surface: stylometric analysis of GPT-4o’s capacity for literary style imitation” ⁴⁹	Directly tackles LLM literary style imitation.	Prompted generation plus distance-based authorship attribution and stylometric comparison.	14 pages (587–600). ⁵⁰	Data availability not obvious in the page snippet I reviewed.	It is methodologically explicit, tied to literary style, and asks a humanities question with computational evidence.
Jan Rybicki, “Can machine translation of literary texts fool stylometry?” ⁵¹	Shows how machine-mediated textual transformation changes stylometric analysis.	Authorship-attribution stylometry plus network visualisation.	9 pages (268–276). ⁵²	Not evident in snippet reviewed.	Very close in spirit to your paper: transformation, literary texts, stylometric measurement, narrow but sharp claim.

DSH article	Contribution	Methods	Approximate length	Data / supplement signal	Why it worked as DSH
Emrah Peksoy, “Who wrote the penultimate chapter of The Female Quixote” ⁵³	Solves a concrete literary-authorship problem.	Bootstrap consensus networks, distance/machine-learning classification, rolling stylometry, cross-validation.	Advance article; page range not yet visible in source I reviewed.	Supplementary Data is explicitly signalled on the article page. ⁵⁴	It is unmistakably a humanities paper first and a methods paper second. That is exactly the balance you need.
G. Tidman and M. Sarazin, “dynamic network visualization and analysis of an historical literary debate” ⁵⁵	Mixed-methods humanities argument built from network analysis.	Social network analysis, controversy mapping, traditional humanities interpretation, reproducible R code.	21 pages (464–484). ⁵⁶	The article text explicitly points readers to reproducible code on GitHub. ⁵⁶	This is a model for how to tie digital method to interpretive payoff and visibly reproducible practice.
Hongwei Zhan, “Key cluster identification in literary texts using and comparing multiple measures” ⁵⁷	Method comparison inside a literary corpus.	Comparative evaluation of five measures on Dickens corpora.	14 pages (668–681). ⁵⁸	Supplementary data is flagged on the issue page. ⁵⁸	It succeeds because it is methodological without ceasing to be textually grounded and corpus-explicit.

The lesson from those exemplars is brutal but simple: **successful DSH papers do not merely run methods on cultural data**. They make the method answer a recognisably humanistic question, document the corpus and assumptions, and remain disciplined about interpretive scope. Your repo already has the documentation. You now need the humanistic question to dominate the manuscript rather than sit behind it. ⁵⁹

Community advice worth taking seriously

The community-level advice is not authoritative in the way OUP and DSH pages are, but it is directionally useful and it lines up with the formal evidence. On Reddit, humanities scholars repeatedly stress that you should **write with the target journal in mind**, that **style matters unusually much in the humanities**, and that review/publication timelines are longer than many STEM newcomers expect. A

DigitalHumanities thread on paper structure also reflects that journal work needs more explicit literature framing than informal conference prose. Practitioner and community resources in stylometry likewise emphasise executable, replicable workflows, while a DH blog essay by James O’Sullivan argues that DH has a reproducibility problem whenever methods, data access, and interpretation are not made transparent enough for peers to evaluate. Those are not passing comments here; they map almost perfectly onto what your repo still needs at manuscript level. ⁶⁰

Open questions and residual limitations

Two uncertainties remain. First, I did not find a clean public statement of **DSH’s exact peer-review model** in the author-facing material I reviewed, so the anonymisation guidance above is a strong precaution rather than a documented journal rule. Second, some article-level details in the exemplar comparison, especially data-availability signals, come from issue pages and article metadata rather than from full-text inspection of every paper, because not all recent DSH content is openly accessible in full. Those uncertainties do not change the main conclusion. They only mean you should verify the final submission mechanics before upload. ⁶¹

The bottom line is blunt: **the repo is good enough to support a serious DSH submission, but not yet good enough to deserve one**. The empirical work is ahead of the manuscript. Fix the framing, close the work/topic loophole, neutralise the warning-row objection, add one visible layer of humanities interpretation, and present the paper as a bounded DH intervention on literary style rather than as a generic AI benchmark. Do that, and the project stops looking like a clever repo and starts looking like an article. ⁶²

¹ ¹⁶ ²⁹ ³⁰ ⁴⁶ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_15_original_to_rewritten_degradation_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_15_original_to_rewritten_degradation_report.md

² ³ ¹⁸ ²¹ ²² <https://academic.oup.com/dsh/pages/About>
<https://academic.oup.com/dsh/pages/About>

⁴ ¹⁵ ²⁸ ³² ³⁶ ⁴¹ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_08_status.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_08_status.md

⁵ ³⁴ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/docs/rewrite_protocol.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/docs/rewrite_protocol.md

⁶ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_13_modeling_matrix_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_13_modeling_matrix_report.md

⁷ ⁴⁰ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/final_reproducibility_audit_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/final_reproducibility_audit_report.md

⁸ https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_14_original_author_baseline_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_14_original_author_baseline_report.md

- 9 10 33 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_10_rewrite_generation_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_10_rewrite_generation_report.md
- 11 14 23 47 62 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/deep_manuscript_readiness_audit.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/deep_manuscript_readiness_audit.md
- 12 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_04_status.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_04_status.md
- 13 <https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/environment.yml>
<https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/environment.yml>
- 17 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_17_burrows_delta_distance_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_17_burrows_delta_distance_report.md
- 19 20 24 25 26 39 43 44 https://academic.oup.com/dsh/pages/General_Instructions
https://academic.oup.com/dsh/pages/General_Instructions
- 27 49 <https://academic.oup.com/dsh/article/40/2/587/8118784>
<https://academic.oup.com/dsh/article/40/2/587/8118784>
- 31 <https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqag040/8586905>
<https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqag040/8586905>
- 35 <https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqag067/8694305>
<https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqag067/8694305>
- 37 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_18_feature_family_vulnerability_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_18_feature_family_vulnerability_report.md
- 38 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_19_tables_figures_report.md
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/logs/step_19_tables_figures_report.md
- 42 https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/.github/workflows/run_step15.yml
https://raw.githubusercontent.com/MuhammadTahaBinZaeem/litpaper/main/.github/workflows/run_step15.yml
- 45 61 **Ethics | Oxford Academic**
[https://academic.oup.com/pages/for-authors/journals/preparing-and-submitting-your-manuscript/ethics?](https://academic.oup.com/pages/for-authors/journals/preparing-and-submitting-your-manuscript/ethics?utm_source=chatgpt.com)
[utm_source=chatgpt.com](https://academic.oup.com/pages/for-authors/journals/preparing-and-submitting-your-manuscript/ethics?utm_source=chatgpt.com)
- 48 https://academic.oup.com/dsh/issue/41/Supplement_1
https://academic.oup.com/dsh/issue/41/Supplement_1
- 50 <https://academic.oup.com/dsh/article-pdf/40/2/587/62995095/fqaf035.pdf>
<https://academic.oup.com/dsh/article-pdf/40/2/587/62995095/fqaf035.pdf>
- 51 <https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqaf010/8024234>
<https://academic.oup.com/dsh/advance-article-abstract/doi/10.1093/llc/fqaf010/8024234>
- 52 **Can machine translation of literary texts fool stylometry?**
https://academic.oup.com/dsh/article-abstract/40/1/268/8024234?utm_source=chatgpt.com

- 53 54 Who wrote the penultimate chapter of The Female Quixote
https://academic.oup.com/dsh/advance-article/doi/10.1093/llc/fqag030/8533199?utm_source=chatgpt.com
- 55 59 dynamic network visualization and analysis of an historical ...
https://academic.oup.com/dsh/article/41/1/464/8424365?utm_source=chatgpt.com
- 56 <https://academic.oup.com/dsh/article/41/1/464/8424365>
<https://academic.oup.com/dsh/article/41/1/464/8424365>
- 57 Key cluster identification in literary texts using and comparing ...
https://academic.oup.com/dsh/article/40/2/668/8124013?utm_source=chatgpt.com
- 58 Volume 40 Issue 2 | Digital Scholarship in the Humanities
https://academic.oup.com/dsh/issue/40/2?utm_source=chatgpt.com
- 60 Humanities: publishing articles : r/academia
https://www.reddit.com/r/academia/comments/168u8ra/humanities_publishing_articles/